

Characterizing Scheduling Heuristics in Low-Bandwidth Links

Tamsin Abernathy, Alina Kowalczyk, Wren Kowalczyk, Elke Jovanovic, Yusuf Kowalczyk

Thornwick College of Engineering

Section: Computing

Published: 2020-06-18

DOI: 10.5281/mjas.9.2.06

Keywords: error propagation; cost accounting; noise robustness; finite elements; benchmarking

Abstract

This paper reports a systematic evaluation of sparse query planners. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Residual error grows approximately linearly with the number of coupled parameters. Calibration drift over the trial period remained below the threshold set out in the protocol.

1. Introduction

Calibration drift over the trial period remained below the threshold set out in the protocol. Neither of the alternative formulations improved on the simple additive model in our setting. The measurements reported here were collected over twelve independent trials under controlled conditions. Taken together, these findings argue for treating the two regimes separately in future analyses.

Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Sensitivity to the smoothing window proved modest across the range we examined. Taken together, these findings argue for treating the two regimes separately in future analyses. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Residual error grows approximately linearly with the number of coupled parameters.

Sensitivity to the smoothing window proved modest across the range we examined. Residual error grows approximately linearly with the number of coupled parameters. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

We release the processing scripts and intermediate artefacts to support independent verification. Residual error grows approximately linearly with the number of coupled parameters. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The reduced model retains most of the explanatory power at a fraction of the computational cost. Sensitivity to the smoothing window proved modest across the range we examined. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Longer observation windows narrow the confidence interval but do not shift the central estimate. Neither of the alternative formulations improved on the simple additive model in our setting. Taken together, these findings argue for treating the two regimes separately in future analyses.

2. Related Work

Residual error grows approximately linearly with the number of coupled parameters. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. The measurements reported here were collected over twelve independent trials under controlled conditions. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Sensitivity to the smoothing window proved modest across the range we examined. Calibration drift over the trial period remained below the threshold set out in the protocol.

Discrepancies at the boundary are consistent with the simplifying assumptions noted above. The measurements reported here were collected over twelve independent trials under controlled conditions. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Residual error grows approximately linearly with the number of coupled parameters.

Longer observation windows narrow the confidence interval but do not shift the central estimate. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Calibration drift over the trial period remained below the threshold set out in the protocol. Neither of the alternative formulations improved on the simple additive model in our setting. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

3. Materials and Methods

Taken together, these findings argue for treating the two regimes separately in future analyses. Calibration drift over the trial period remained below the threshold set out in the protocol. Sensitivity to the smoothing window proved modest across the range we examined. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We release the processing scripts and intermediate artefacts to support independent verification. The measurements reported here were collected over twelve independent trials under controlled conditions.

The reduced model retains most of the explanatory power at a fraction of the computational cost. Calibration drift over the trial period remained below the threshold set out in the protocol. Neither of the alternative formulations improved on the simple additive model in our setting. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

The measurements reported here were collected over twelve independent trials under controlled conditions. Sensitivity to the smoothing window proved modest across the range we examined. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Residual error grows approximately linearly with the number of coupled parameters. Longer observation windows narrow the confidence interval but do not shift the central estimate. The reduced model retains most of the explanatory power at a fraction of the computational cost. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance.

4. Experimental Setup

We release the processing scripts and intermediate artefacts to support independent verification. The reduced model retains most of the explanatory power at a fraction of the computational cost. Longer observation windows narrow the confidence interval but do not shift the central estimate. Taken together, these findings argue for treating the two regimes separately in future analyses. Calibration drift over the trial period remained below the threshold set out in the protocol. Sensitivity to the smoothing window proved modest across the range we examined. The measurements reported here were collected over twelve independent trials under controlled conditions. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Calibration drift over the trial period remained below the threshold set out in the protocol. The measurements reported here were collected over twelve independent trials under controlled conditions. The reduced model retains most of the explanatory power at a fraction of the computational cost. Neither of the alternative formulations improved on the simple additive model in our setting. Taken together, these findings argue for treating the two regimes separately in future analyses. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Sensitivity to the smoothing window proved modest across the range we examined.

Calibration drift over the trial period remained below the threshold set out in the protocol. Sensitivity to the smoothing window proved modest across the range we examined. We release the processing scripts and intermediate artefacts to support independent verification. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Taken together, these findings argue for treating the two regimes separately in future analyses. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

5. Results

Sensitivity to the smoothing window proved modest across the range we examined. The reduced model retains most of the explanatory power at a fraction of the computational cost. Calibration drift over the trial period remained below the threshold set out in the protocol. We release the processing scripts and intermediate artefacts to support independent verification.

Neither of the alternative formulations improved on the simple additive model in our setting. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Sensitivity to the smoothing window proved modest across the range we examined. Longer observation windows narrow the confidence interval but do not shift the central estimate. Residual error grows approximately linearly with the number of coupled parameters.

We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Sensitivity to the smoothing window proved modest across the range we examined. Residual error grows approximately linearly with the number of coupled parameters. Longer observation windows narrow the confidence interval but do not shift the central estimate. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Neither of the alternative formulations improved on the simple additive model in our setting. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

Sensitivity to the smoothing window proved modest across the range we examined. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The measurements reported here were collected over twelve independent trials under controlled conditions. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

6. Discussion

Calibration drift over the trial period remained below the threshold set out in the protocol. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Neither of the alternative formulations improved on the simple additive model in our setting. Residual error grows approximately linearly with the number of coupled parameters. Longer observation windows narrow the confidence interval but do not shift the central estimate. We release the processing scripts and intermediate artefacts to support independent verification. Sensitivity to the smoothing window proved modest across the range we examined. Taken together, these findings argue for treating the two regimes separately in future analyses.

Neither of the alternative formulations improved on the simple additive model in our setting. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Longer observation windows narrow the confidence interval but do not shift the central estimate. We release the processing scripts and intermediate artefacts to support independent verification. Sensitivity to the smoothing window proved modest across the range we examined.

The reduced model retains most of the explanatory power at a fraction of the computational cost. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Calibration drift over the trial period remained below the threshold set out in the protocol.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Longer observation windows narrow the confidence interval but do not shift the central estimate. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

7. Limitations

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The reduced model retains most of the explanatory power at a fraction of the computational cost. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Residual error grows approximately linearly with the number of coupled parameters. Calibration drift over the trial period remained below the threshold set out in the protocol.

Neither of the alternative formulations improved on the simple additive model in our setting. Calibration drift over the trial period remained below the threshold set out in the protocol. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Longer observation windows narrow the confidence interval but do not shift the central estimate.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. Residual error

grows approximately linearly with the number of coupled parameters. Longer observation windows narrow the confidence interval but do not shift the central estimate. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Taken together, these findings argue for treating the two regimes separately in future analyses. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. The measurements reported here were collected over twelve independent trials under controlled conditions. Neither of the alternative formulations improved on the simple additive model in our setting.

Residual error grows approximately linearly with the number of coupled parameters. We release the processing scripts and intermediate artefacts to support independent verification. The reduced model retains most of the explanatory power at a fraction of the computational cost. Taken together, these findings argue for treating the two regimes separately in future analyses. Sensitivity to the smoothing window proved modest across the range we examined. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

8. Conclusion

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Sensitivity to the smoothing window proved modest across the range we examined. We release the processing scripts and intermediate artefacts to support independent verification. The reduced model retains most of the explanatory power at a fraction of the computational cost.

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Sensitivity to the smoothing window proved modest across the range we examined. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Longer observation windows narrow the confidence interval but do not shift the central estimate. We release the processing scripts and intermediate artefacts to support independent verification. Taken together, these findings argue for treating the two regimes separately in future analyses.

The measurements reported here were collected over twelve independent trials under controlled conditions. Calibration drift over the trial period remained below the threshold set out in the protocol. Longer observation windows narrow the confidence interval but do not shift the central estimate. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. The reduced model retains most of the explanatory power at a fraction of the computational cost. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Taken together, these findings argue for treating the two regimes separately in future analyses.