

Toward Sparse Scheduling Heuristics: Evidence from Multi-Tenant Pipelines

Cordelia Nakashima

Calder Bay University

Section: Computing
Published: 2020-12-16
DOI: 10.5281/mjas.9.4.03
Keywords: ablation study; cost accounting; cross-validation

Abstract

Work in computing research has long relied on assumptions we test directly here. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Sensitivity to the smoothing window proved modest across the range we examined. Taken together, these findings argue for treating the two regimes separately in future analyses.

1. Introduction

Neither of the alternative formulations improved on the simple additive model in our setting. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The measurements reported here were collected over twelve independent trials under controlled conditions. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

Sensitivity to the smoothing window proved modest across the range we examined. Residual error grows approximately linearly with the number of coupled parameters. Calibration drift over the trial period remained below the threshold set out in the protocol. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Longer observation windows narrow the confidence interval but do not shift the central estimate. We release the processing scripts and intermediate artefacts to support independent verification. Neither of the alternative formulations improved on the simple additive model in our setting. The measurements reported here were collected over twelve independent trials under controlled conditions. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

Calibration drift over the trial period remained below the threshold set out in the protocol. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Taken together, these findings argue for treating the two regimes separately in future analyses. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Neither of the alternative formulations improved on the simple additive model in our setting.

2. Related Work

Residual error grows approximately linearly with the number of coupled parameters. Calibration drift over the trial period remained below the threshold set out in the protocol. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Longer observation windows narrow the confidence interval but do not shift the central estimate. We release the processing scripts and intermediate artefacts to support independent verification.

We release the processing scripts and intermediate artefacts to support independent verification. The measurements reported here were collected over twelve independent trials under controlled conditions. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Taken together, these findings argue for treating the two regimes separately in future analyses.

Taken together, these findings argue for treating the two regimes separately in future analyses. Neither of the alternative formulations improved on the simple additive model in our setting. We release the processing scripts and intermediate artefacts to support independent verification. Longer observation windows narrow the confidence interval but do not shift the central estimate. Sensitivity to the smoothing window proved modest across the range we examined. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The reduced model retains most of the explanatory power at a fraction of the computational cost. The measurements reported here were collected over twelve independent trials under controlled conditions.

3. Materials and Methods

We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. We release the processing scripts and intermediate artefacts to support independent verification. Neither of the alternative formulations improved on the simple additive model in our setting. Residual error grows approximately linearly with the number of coupled parameters.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Calibration drift over the trial period remained below the threshold set out in the protocol. Residual error grows approximately linearly with the number of coupled parameters. The reduced model retains most of the explanatory power at a fraction of the computational cost. Longer observation windows narrow the confidence interval but do not shift the central estimate.

4. Experimental Setup

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. We release the processing scripts and intermediate artefacts to support independent verification. Residual error grows approximately linearly with the number of coupled parameters. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

Sensitivity to the smoothing window proved modest across the range we examined. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Taken together, these findings argue for

treating the two regimes separately in future analyses. Replicate agreement was high, suggesting the protocol is robust to small handling differences. We release the processing scripts and intermediate artefacts to support independent verification. Neither of the alternative formulations improved on the simple additive model in our setting.

Calibration drift over the trial period remained below the threshold set out in the protocol. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Residual error grows approximately linearly with the number of coupled parameters. The reduced model retains most of the explanatory power at a fraction of the computational cost. We release the processing scripts and intermediate artefacts to support independent verification. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The measurements reported here were collected over twelve independent trials under controlled conditions.

5. Results

Longer observation windows narrow the confidence interval but do not shift the central estimate. The reduced model retains most of the explanatory power at a fraction of the computational cost. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Taken together, these findings argue for treating the two regimes separately in future analyses. The measurements reported here were collected over twelve independent trials under controlled conditions.

Discrepancies at the boundary are consistent with the simplifying assumptions noted above. The reduced model retains most of the explanatory power at a fraction of the computational cost. Residual error grows approximately linearly with the number of coupled parameters. Neither of the alternative formulations improved on the simple additive model in our setting. Sensitivity to the smoothing window proved modest across the range we examined.

6. Discussion

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We release the processing scripts and intermediate artefacts to support independent verification. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The reduced model retains most of the explanatory power at a fraction of the computational cost. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. The measurements reported here were collected over twelve independent trials under controlled conditions. Longer observation windows narrow the confidence interval but do not shift the central estimate.

Calibration drift over the trial period remained below the threshold set out in the protocol. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The measurements reported here were collected over twelve independent trials under controlled conditions. We release the processing scripts and intermediate artefacts to support independent verification. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Sensitivity to the smoothing window proved modest across the range we examined. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Taken together, these findings argue for treating the two regimes separately in future analyses.

We release the processing scripts and intermediate artefacts to support independent verification. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Calibration drift over the trial period remained below the threshold set out in the protocol. The reduced model retains most of the explanatory power at a fraction of the computational cost. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Residual error grows approximately linearly with the number of coupled parameters.

The reduced model retains most of the explanatory power at a fraction of the computational cost. We release the processing scripts and intermediate artefacts to support independent verification. Taken together, these findings argue for treating the two regimes separately in future analyses. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity.

7. Limitations

The measurements reported here were collected over twelve independent trials under controlled conditions. We release the processing scripts and intermediate artefacts to support independent verification. Sensitivity to the smoothing window proved modest across the range we examined. Neither of the alternative formulations improved on the simple additive model in our setting. Longer observation windows narrow the confidence interval but do not shift the central estimate.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. Sensitivity to the smoothing window proved modest across the range we examined. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity.

Longer observation windows narrow the confidence interval but do not shift the central estimate. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Sensitivity to the smoothing window proved modest across the range we examined. We release the processing scripts and intermediate artefacts to support independent verification. The measurements reported here were collected over twelve independent trials under controlled conditions. Calibration drift over the trial period remained below the threshold set out in the protocol. Neither of the alternative formulations improved on the simple additive model in our setting. Taken together, these findings argue for treating the two regimes separately in future analyses.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Taken together, these findings argue for treating the two regimes separately in future analyses. The reduced model retains most of the explanatory power at a fraction of the computational cost. Longer observation windows narrow the confidence interval but do not shift the central estimate. Sensitivity to the smoothing window proved modest across the range we examined. Calibration drift over the trial period remained below the threshold set out in the protocol.

Neither of the alternative formulations improved on the simple additive model in our setting. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. The measurements reported here were collected over twelve independent trials under controlled conditions. Longer observation windows narrow the confidence interval but do not shift the central estimate. Calibration drift over the trial period remained below the threshold set out in the protocol. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The reduced model retains most of the explanatory power at a fraction of the computational cost. Taken together, these findings argue for treating the two regimes

separately in future analyses.