

A Comparative Study of Hierarchical Query Planners

Priya Whitlock, Jarrah Nakashima, Devraj Fontenot, Eero Yarborough

Northfield Polytechnic

Section: Computing
Published: 2021-06-21
DOI: 10.5281/mjas.10.2.05
Keywords: sensitivity analysis; calibration; cohort design

Abstract

We examine the behaviour of workload forecasts in append-only stores. Sensitivity to the smoothing window proved modest across the range we examined. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Neither of the alternative formulations improved on the simple additive model in our setting.

1. Introduction

We release the processing scripts and intermediate artefacts to support independent verification. Calibration drift over the trial period remained below the threshold

artefacts to support independent verification. The measurements reported here were collected over twelve independent trials under controlled conditions.

We release the processing scripts and intermediate artefacts to support independent verification. Taken together, these findings argue for treating the two regimes separately in future analyses. The measurements reported here were collected over twelve independent trials under controlled conditions. Calibration drift over the trial period remained below the threshold set out in the protocol. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Neither of the alternative formulations improved on the simple additive model in our setting.

Residual error grows approximately linearly with the number of coupled parameters. Sensitivity to the smoothing window proved modest across the range we examined. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Taken together, these findings argue for treating the two regimes separately in future analyses. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. The measurements reported here were collected over twelve independent trials under controlled conditions. The reduced model retains most of the explanatory power at a fraction of the computational cost.

3. Materials and Methods

The measurements reported here were collected over twelve independent trials under controlled conditions. We release the processing scripts and intermediate artefacts to support independent verification. Neither of the alternative formulations improved on the simple additive model in our setting. Longer observation windows narrow the confidence

and intermediate artefacts to support independent verification. The measurements reported here were collected over twelve independent trials under controlled conditions.

Calibration drift over the trial period remained below the threshold set out in the protocol. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The reduced model retains most of the explanatory power at a fraction of the computational cost. The measurements reported here were collected over twelve independent trials under controlled conditions. Residual error grows approximately linearly with the number of coupled parameters. Neither of the alternative formulations improved on the simple additive model in our setting.

We release the processing scripts and intermediate artefacts to support independent verification. Taken together, these findings argue for treating the two regimes separately in future analyses. The reduced model retains most of the explanatory power at a fraction of the computational cost. Longer observation windows narrow the confidence interval but do not shift the central estimate. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

Discrepancies at the boundary are consistent with the simplifying assumptions noted above. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We release the processing scripts and intermediate artefacts to support independent verification. Sensitivity to the smoothing window proved modest across the range we examined. The reduced model retains most of the explanatory power at a fraction of the computational cost. Taken together, these findings argue for treating the two regimes separately in future analyses. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Calibration drift over the trial period remained below the threshold set out in the protocol. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Longer observation windows narrow the confidence interval but do not shift the central estimate. Neither of the alternative formulations improved on the simple additive model in our setting. The reduced model retains most of the explanatory power at a fraction of the computational cost.

6. Discussion

The measurements reported here were collected over twelve independent trials under controlled conditions. Calibration drift over the trial period remained below the threshold set out in the protocol. We release the processing scripts and intermediate artefacts to support independent verification. Neither of the alternative formulations improved on the simple additive model in our setting. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. The reduced model retains most of the explanatory power at a fraction of the computational cost. Taken together, these findings argue for treating the two regimes separately in future analyses. We release the processing scripts and intermediate artefacts to support independent verification. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Residual error grows approximately linearly with the number of coupled parameters. Longer observation windows narrow the confidence interval but do not shift the central estimate. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Neither of the alternative formulations improved on the simple additive model in our setting.