

Toward Hierarchical Index Structures: Evidence from Append-Only Stores

Imogen Xanthos, Alina Kowalczyk, Fenna Strand

Institute for Applied Dynamics, Varnholt

Section: Computing

Published: 2022-06-24

DOI: 10.5281/mjas.11.2.04

Keywords: ablation study; parameter sweep; calibration; reproducibility; error propagation

Abstract

Work in computing research has long relied on assumptions we test directly here. Taken together, these findings argue for treating the two regimes separately in future analyses. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Longer observation windows narrow the confidence interval but do not shift the central estimate.

1. Introduction

The measurements reported here were collected over twelve independent trials under controlled conditions. Taken together, these findings argue for treating the two regimes separately in future analyses. Residual error grows approximately linearly with the number of coupled parameters. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

The reduced model retains most of the explanatory power at a fraction of the computational cost. Calibration drift over the trial period remained below the threshold set out in the protocol. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Residual error grows approximately linearly with the number of coupled parameters. Neither of the alternative formulations improved on the simple additive model in our setting. Sensitivity to the smoothing window proved modest across the range we examined. We release the processing scripts and intermediate artefacts to support independent verification.

Neither of the alternative formulations improved on the simple additive model in our setting. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Residual error grows approximately linearly with the number of coupled parameters. Sensitivity to the smoothing window proved modest across the range we examined. Taken together, these findings argue for treating the two regimes separately in future analyses. We release the processing scripts and intermediate artefacts to support independent verification. The reduced model retains most of the explanatory power at a fraction of the computational cost.

2. Related Work

The measurements reported here were collected over twelve independent trials under controlled conditions. Calibration drift over the trial period remained below the threshold set out in the protocol. Neither of the alternative formulations improved on the simple additive model in our setting. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Residual error grows approximately linearly with the number of coupled parameters. Longer observation windows narrow the confidence interval but do not shift the central estimate. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity.

The reduced model retains most of the explanatory power at a fraction of the computational cost. Neither of the alternative formulations improved on the simple additive model in our setting. The measurements reported here were collected over twelve independent trials under controlled conditions. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Longer observation windows narrow the confidence interval but do not shift the central estimate. We release the processing scripts and intermediate artefacts to support independent verification.

3. Materials and Methods

Replicate agreement was high, suggesting the protocol is robust to small handling differences. Calibration drift over the trial period remained below the threshold set out in the protocol. Neither of the alternative formulations improved on the simple additive model in our setting. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Longer observation windows narrow the confidence interval but do not shift the central estimate. We release the processing scripts and intermediate artefacts to support independent verification.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Residual error grows approximately linearly with the number of coupled parameters. Neither of the alternative formulations improved on the simple additive model in our setting. Taken together, these findings argue for treating the two regimes separately in future analyses.

Sensitivity to the smoothing window proved modest across the range we examined. We release the processing scripts and intermediate artefacts to support independent verification. The reduced model retains most of the explanatory power at a fraction of the computational cost. The measurements reported here were collected over twelve independent trials under controlled conditions. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Longer observation windows narrow the confidence interval but do not shift the central estimate. Calibration drift over the trial period remained below the threshold set out in the protocol. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

We release the processing scripts and intermediate artefacts to support independent verification. Neither of the alternative formulations improved on the simple additive model in our setting. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Longer observation windows narrow the confidence interval but do not shift the central estimate.

4. Experimental Setup

Residual error grows approximately linearly with the number of coupled parameters. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Neither of the alternative formulations improved on the simple additive model in our setting. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Taken together, these findings argue for treating the two regimes separately in future analyses. The measurements reported here were collected over twelve independent trials under controlled conditions.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Calibration drift over the trial period remained below the threshold set out in the protocol. Longer observation windows narrow the confidence interval but do not shift the central estimate. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We release the processing scripts and intermediate artefacts to support independent verification. The reduced model retains most of the explanatory power at a fraction of the computational cost.

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Residual error grows approximately linearly with the number of coupled parameters.

5. Results

Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Taken together, these findings argue for treating the two regimes separately in future analyses. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Residual error grows approximately linearly with the number of coupled parameters. The reduced model retains most of the explanatory power at a fraction of the computational cost.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We release the processing scripts and intermediate artefacts to support independent verification. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Neither of the alternative formulations improved on the simple additive model in our setting. The measurements reported here were collected over twelve independent trials under controlled conditions. Longer observation windows narrow the confidence interval but do not shift the central estimate.

The measurements reported here were collected over twelve independent trials under controlled conditions. Longer observation windows narrow the confidence interval but do not shift the central estimate. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. We release the processing scripts and intermediate artefacts to support independent verification. Residual error grows approximately linearly with the number of coupled parameters. Taken together, these findings argue for treating the two regimes

separately in future analyses.