

Probabilistic Query Planners for Edge Deployments

Kenji Nakashima, Zofia Ueckermann

Sableport Marine Laboratory

Section: Computing

Published: 2022-06-28

DOI: 10.5281/mjas.11.2.06

Keywords: ablation study; surrogate modelling; cross-validation; noise robustness

Abstract

We examine the behaviour of graph embeddings in heterogeneous clusters. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Longer observation windows narrow the confidence interval but do not shift the central estimate. We release the processing scripts and intermediate artefacts to support independent verification.

1. Introduction

We release the processing scripts and intermediate artefacts to support independent verification. Sensitivity to the smoothing window proved modest across the range we examined. Calibration drift over the trial period remained below the threshold set out in the protocol. Residual error grows approximately linearly with the number of coupled parameters. Taken together, these findings argue for treating the two regimes separately in future analyses. The reduced model retains most of the explanatory power at a fraction of the computational cost. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. The measurements reported here were collected over twelve independent trials under controlled conditions. Neither of the alternative formulations improved on the simple additive model in our setting. Residual error grows approximately linearly with the number of coupled parameters. Longer observation windows narrow the confidence interval but do not shift the central estimate. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Calibration drift over the trial period remained below the threshold set out in the protocol. Taken together, these findings argue for treating the two regimes separately in future analyses.

We release the processing scripts and intermediate artefacts to support independent verification. Taken together, these findings argue for treating the two regimes separately in future analyses. Neither of the alternative formulations improved on the simple additive model in our setting. Sensitivity to the smoothing window proved modest across the range we examined. Calibration drift over the trial period remained below the threshold set out in the protocol.

2. Related Work

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Sensitivity to the smoothing window proved modest across the range we examined. Longer observation windows narrow the confidence interval but do not shift the central estimate. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Residual error grows approximately linearly with the number of coupled parameters.

Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Taken together, these findings argue for treating the two regimes separately in future analyses. The reduced model retains most of the explanatory power at a fraction of the computational cost. Sensitivity to the smoothing window proved modest across the range we examined. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Longer observation windows narrow the confidence interval but do not shift the central estimate. Residual error grows approximately linearly with the number of coupled parameters. Neither of the alternative formulations improved on the simple additive model in our setting.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We release the processing scripts and intermediate artefacts to support independent verification. Taken together, these findings argue for treating the two regimes separately in future analyses. Calibration drift over the trial period remained below the threshold set out in the protocol. Residual error grows approximately linearly with the number of coupled parameters. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Sensitivity to the smoothing window proved modest across the range we examined. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

Sensitivity to the smoothing window proved modest across the range we examined. The measurements reported here were collected over twelve independent trials under controlled conditions. We release the processing scripts and intermediate artefacts to support independent verification. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Residual error grows approximately linearly with the number of coupled parameters. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Longer observation windows narrow the confidence interval but do not shift the central estimate. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity.

We release the processing scripts and intermediate artefacts to support independent verification. Neither of the alternative formulations improved on the simple additive model in our setting. The measurements reported here were collected over twelve independent trials under controlled conditions. Residual error grows approximately linearly with the number of coupled parameters. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Longer observation windows narrow the confidence interval but do not shift the central estimate. Sensitivity to the smoothing window proved modest across the range we examined.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Residual error grows approximately linearly with the number of coupled parameters. Sensitivity to the smoothing window proved modest across the range we examined. Longer observation windows narrow the confidence interval but do not shift the central estimate. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Calibration drift over the trial period remained below the threshold set out in the protocol. The

measurements reported here were collected over twelve independent trials under controlled conditions.

3. Materials and Methods

Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Neither of the alternative formulations improved on the simple additive model in our setting. Sensitivity to the smoothing window proved modest across the range we examined. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The measurements reported here were collected over twelve independent trials under controlled conditions.

Neither of the alternative formulations improved on the simple additive model in our setting. Residual error grows approximately linearly with the number of coupled parameters. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. We release the processing scripts and intermediate artefacts to support independent verification. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. The measurements reported here were collected over twelve independent trials under controlled conditions. Calibration drift over the trial period remained below the threshold set out in the protocol.

Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Sensitivity to the smoothing window proved modest across the range we examined. We release the processing scripts and intermediate artefacts to support independent verification. Neither of the alternative formulations improved on the simple additive model in our setting. Taken together, these findings argue for treating the two regimes separately in future analyses. The measurements reported here were collected over twelve independent trials under controlled conditions. The reduced model retains most of the explanatory power at a fraction of the computational cost. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Calibration drift over the trial period remained below the threshold set out in the protocol.

Discrepancies at the boundary are consistent with the simplifying assumptions noted above. The reduced model retains most of the explanatory power at a fraction of the computational cost. Longer observation windows narrow the confidence interval but do not shift the central estimate. Calibration drift over the trial period remained below the threshold set out in the protocol.

4. Experimental Setup

Taken together, these findings argue for treating the two regimes separately in future analyses. Neither of the alternative formulations improved on the simple additive model in our setting. Longer observation windows narrow the confidence interval but do not shift the central estimate. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Sensitivity to the smoothing window proved modest across the range we examined. The reduced model retains most of the explanatory power at a fraction of the computational cost. Calibration drift over the trial period remained below the threshold set out in the protocol.

Neither of the alternative formulations improved on the simple additive model in our setting. Residual error grows approximately linearly with the number of coupled parameters. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Sensitivity to the smoothing window proved modest

across the range we examined. Longer observation windows narrow the confidence interval but do not shift the central estimate. We release the processing scripts and intermediate artefacts to support independent verification.

5. Results

Taken together, these findings argue for treating the two regimes separately in future analyses. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Residual error grows approximately linearly with the number of coupled parameters. Calibration drift over the trial period remained below the threshold set out in the protocol. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. We release the processing scripts and intermediate artefacts to support independent verification. Residual error grows approximately linearly with the number of coupled parameters. Sensitivity to the smoothing window proved modest across the range we examined. Longer observation windows narrow the confidence interval but do not shift the central estimate. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Taken together, these findings argue for treating the two regimes separately in future analyses. Calibration drift over the trial period remained below the threshold set out in the protocol.

Neither of the alternative formulations improved on the simple additive model in our setting. The measurements reported here were collected over twelve independent trials under controlled conditions. Taken together, these findings argue for treating the two regimes separately in future analyses. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Sensitivity to the smoothing window proved modest across the range we examined. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Longer observation windows narrow the confidence interval but do not shift the central estimate.

Residual error grows approximately linearly with the number of coupled parameters. The reduced model retains most of the explanatory power at a fraction of the computational cost. We release the processing scripts and intermediate artefacts to support independent verification. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. The measurements reported here were collected over twelve independent trials under controlled conditions.

6. Discussion

Taken together, these findings argue for treating the two regimes separately in future analyses. Residual error grows approximately linearly with the number of coupled parameters. Neither of the alternative formulations improved on the simple additive model in our setting. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The reduced model retains most of the explanatory power at a fraction of the computational cost. Sensitivity to the smoothing window proved modest across the range we examined. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Residual error grows approximately linearly with the number of coupled parameters. Discrepancies at the boundary are consistent

with the simplifying assumptions noted above. We release the processing scripts and intermediate artefacts to support independent verification. Neither of the alternative formulations improved on the simple additive model in our setting.

We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Longer observation windows narrow the confidence interval but do not shift the central estimate. Sensitivity to the smoothing window proved modest across the range we examined. The measurements reported here were collected over twelve independent trials under controlled conditions.