

Probabilistic Query Planners for Serverless Runtimes

Sten Vasquez, Jarrah Cisneros, Cordelia Gorecki

Meridian Institute of Technology

Section: Computing

Published: 2023-09-17

DOI: 10.5281/mjas.12.3.04

Keywords: sensitivity analysis; benchmarking; scaling laws; finite elements

Abstract

We examine the behaviour of scheduling heuristics in append-only stores. We release the processing scripts and intermediate artefacts to support independent verification. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Residual error grows approximately linearly with the number of coupled parameters.

1. Introduction

Residual error grows approximately linearly with the number of coupled parameters. Sensitivity to the smoothing window proved modest across the range we examined. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The measurements reported here were collected over twelve independent trials under controlled conditions. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Taken together, these findings argue for treating the two regimes separately in future analyses. Calibration drift over the trial period remained below the threshold set out in the protocol. The reduced model retains most of the explanatory power at a fraction of the computational cost.

The measurements reported here were collected over twelve independent trials under controlled conditions. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Neither of the alternative formulations improved on the simple additive model in our setting. We release the processing scripts and intermediate artefacts to support independent verification. Calibration drift over the trial period remained below the threshold set out in the protocol. Longer observation windows narrow the confidence interval but do not shift the central estimate. The reduced model retains most of the explanatory power at a fraction of the computational cost. Sensitivity to the smoothing window proved modest across the range we examined.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. Residual error grows approximately linearly with the number of coupled parameters. Taken together, these findings argue for treating the two regimes separately in future analyses. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. The reduced model retains most of the explanatory power at a fraction of the computational cost. Neither of the alternative formulations improved on the simple additive model in our setting. We release the processing scripts and intermediate artefacts to support independent verification.

We release the processing scripts and intermediate artefacts to support independent verification. The

measurements reported here were collected over twelve independent trials under controlled conditions. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Longer observation windows narrow the confidence interval but do not shift the central estimate.

2. Related Work

Taken together, these findings argue for treating the two regimes separately in future analyses. Longer observation windows narrow the confidence interval but do not shift the central estimate. Neither of the alternative formulations improved on the simple additive model in our setting. Calibration drift over the trial period remained below the threshold set out in the protocol.

The measurements reported here were collected over twelve independent trials under controlled conditions. The reduced model retains most of the explanatory power at a fraction of the computational cost. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Residual error grows approximately linearly with the number of coupled parameters. Neither of the alternative formulations improved on the simple additive model in our setting. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Taken together, these findings argue for treating the two regimes separately in future analyses. The reduced model retains most of the explanatory power at a fraction of the computational cost. Neither of the alternative formulations improved on the simple additive model in our setting. Replicate agreement was high, suggesting the protocol is robust to small handling differences. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Sensitivity to the smoothing window proved modest across the range we examined. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The measurements reported here were collected over twelve independent trials under controlled conditions.

3. Materials and Methods

The reduced model retains most of the explanatory power at a fraction of the computational cost. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. The measurements reported here were collected over twelve independent trials under controlled conditions. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Sensitivity to the smoothing window proved modest across the range we examined.

The reduced model retains most of the explanatory power at a fraction of the computational cost. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We release the processing scripts and intermediate artefacts to support independent verification. Longer observation windows narrow the confidence interval but do not shift the central estimate. The measurements reported here were collected over twelve independent trials under controlled conditions.

The reduced model retains most of the explanatory power at a fraction of the computational cost. We release the processing scripts and intermediate artefacts to support independent verification. Sensitivity to the smoothing window proved modest across the range we examined. The measurements reported here were collected over twelve independent trials under controlled conditions. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Neither of the alternative formulations improved on the

simple additive model in our setting.

4. Experimental Setup

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Taken together, these findings argue for treating the two regimes separately in future analyses. Calibration drift over the trial period remained below the threshold set out in the protocol. The measurements reported here were collected over twelve independent trials under controlled conditions. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Residual error grows approximately linearly with the number of coupled parameters. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance.

Calibration drift over the trial period remained below the threshold set out in the protocol. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. The measurements reported here were collected over twelve independent trials under controlled conditions. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Taken together, these findings argue for treating the two regimes separately in future analyses. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity.

The measurements reported here were collected over twelve independent trials under controlled conditions. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Taken together, these findings argue for treating the two regimes separately in future analyses. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

5. Results

Calibration drift over the trial period remained below the threshold set out in the protocol. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The reduced model retains most of the explanatory power at a fraction of the computational cost. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The measurements reported here were collected over twelve independent trials under controlled conditions. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity.

Neither of the alternative formulations improved on the simple additive model in our setting. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The measurements reported here were collected over twelve independent trials under controlled conditions. Sensitivity to the smoothing window proved modest across the range we examined. The reduced model retains most of the explanatory power at a fraction of the computational cost. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance.

Calibration drift over the trial period remained below the threshold set out in the protocol. The reduced model retains most of the explanatory power at a fraction of the computational cost. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Longer observation windows narrow the confidence interval but do not shift the central estimate. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Sensitivity to the smoothing window proved modest across the range we examined. Taken together, these findings argue for treating the two regimes separately in future analyses.

Neither of the alternative formulations improved on the simple additive model in our setting. We release the processing scripts and intermediate artefacts to support independent verification. The reduced model retains most of the explanatory power at a fraction of the computational cost. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Calibration drift over the trial period remained below the threshold set out in the protocol. Longer observation windows narrow the confidence interval but do not shift the central estimate. Taken together, these findings argue for treating the two regimes separately in future analyses.

6. Discussion

We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Taken together, these findings argue for treating the two regimes separately in future analyses. The reduced model retains most of the explanatory power at a fraction of the computational cost. Residual error grows approximately linearly with the number of coupled parameters. We release the processing scripts and intermediate artefacts to support independent verification. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The measurements reported here were collected over twelve independent trials under controlled conditions. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

Taken together, these findings argue for treating the two regimes separately in future analyses. The measurements reported here were collected over twelve independent trials under controlled conditions. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The reduced model retains most of the explanatory power at a fraction of the computational cost. Residual error grows approximately linearly with the number of coupled parameters. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Calibration drift over the trial period remained below the threshold set out in the protocol.

7. Limitations

Neither of the alternative formulations improved on the simple additive model in our setting. Residual error grows approximately linearly with the number of coupled parameters. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The measurements reported here were collected over twelve independent trials under controlled conditions.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Neither of the alternative formulations improved on the simple additive model in our setting. The reduced model retains most of the explanatory power at a fraction of the computational cost. Longer observation windows narrow the confidence interval but do not shift the central estimate. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Sensitivity to the smoothing window proved modest across the range we

examined. We release the processing scripts and intermediate artefacts to support independent verification.

Residual error grows approximately linearly with the number of coupled parameters. The measurements reported here were collected over twelve independent trials under controlled conditions. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Calibration drift over the trial period remained below the threshold set out in the protocol. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Taken together, these findings argue for treating the two regimes separately in future analyses.

Residual error grows approximately linearly with the number of coupled parameters. Replicate agreement was high, suggesting the protocol is robust to small handling differences. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Taken together, these findings argue for treating the two regimes separately in future analyses.

8. Conclusion

The reduced model retains most of the explanatory power at a fraction of the computational cost. The measurements reported here were collected over twelve independent trials under controlled conditions. Neither of the alternative formulations improved on the simple additive model in our setting. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Replicate agreement was high, suggesting the protocol is robust to small handling differences. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Sensitivity to the smoothing window proved modest across the range we examined. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

Longer observation windows narrow the confidence interval but do not shift the central estimate. The reduced model retains most of the explanatory power at a fraction of the computational cost. Neither of the alternative formulations improved on the simple additive model in our setting. Taken together, these findings argue for treating the two regimes separately in future analyses. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity.

The measurements reported here were collected over twelve independent trials under controlled conditions. Sensitivity to the smoothing window proved modest across the range we examined. Calibration drift over the trial period remained below the threshold set out in the protocol. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Longer observation windows narrow the confidence interval but do not shift the central estimate. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Replicate agreement was high, suggesting the protocol is robust to small handling differences. We release the processing scripts and intermediate artefacts to support independent verification.