

Toward Adaptive Cache Hierarchies: Evidence from Serverless Runtimes

Hyun-woo Cisneros, Mathis Nakashima

Institute for Applied Dynamics, Varnholt

Section: Computing

Published: 2023-12-22

DOI: 10.5281/mjas.12.4.02

Keywords: reproducibility; cost accounting; cohort design; finite elements

Abstract

We examine the behaviour of streaming joins in low-bandwidth links. Longer observation windows narrow the confidence interval but do not shift the central estimate. Residual error grows approximately linearly with the number of coupled parameters. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

1. Introduction

Replicate agreement was high, suggesting the protocol is robust to small handling differences. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Longer observation windows narrow the confidence interval but do not shift the central estimate. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The measurements reported here were collected over twelve independent trials under controlled conditions. Residual error grows approximately linearly with the number of coupled parameters.

We release the processing scripts and intermediate artefacts to support independent verification. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Neither of the alternative formulations improved on the simple additive model in our setting. The measurements reported here were collected over twelve independent trials under controlled conditions. The reduced model retains most of the explanatory power at a fraction of the computational cost. Taken together, these findings argue for treating the two regimes separately in future analyses. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity.

Sensitivity to the smoothing window proved modest across the range we examined. The reduced model retains most of the explanatory power at a fraction of the computational cost. Longer observation windows narrow the confidence interval but do not shift the central estimate. We release the processing scripts and intermediate artefacts to support independent verification. Calibration drift over the trial period remained below the threshold set out in the protocol. Neither of the alternative formulations improved on the simple additive model

in our setting. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance.

2. Related Work

The measurements reported here were collected over twelve independent trials under controlled conditions. Calibration drift over the trial period remained below the threshold set out in the protocol. Taken together, these findings argue for treating the two regimes separately in future analyses. Neither of the alternative formulations improved on the simple additive model in our setting. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

Longer observation windows narrow the confidence interval but do not shift the central estimate. Taken together, these findings argue for treating the two regimes separately in future analyses. Neither of the alternative formulations improved on the simple additive model in our setting. Calibration drift over the trial period remained below the threshold set out in the protocol. The reduced model retains most of the explanatory power at a fraction of the computational cost. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We release the processing scripts and intermediate artefacts to support independent verification. Residual error grows approximately linearly with the number of coupled parameters. The measurements reported here were collected over twelve independent trials under controlled conditions.

Taken together, these findings argue for treating the two regimes separately in future analyses. Longer observation windows narrow the confidence interval but do not shift the central estimate. The measurements reported here were collected over twelve independent trials under controlled conditions. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Calibration drift over the trial period remained below the threshold set out in the protocol. Neither of the alternative formulations improved on the simple additive model in our setting. The reduced model retains most of the explanatory power at a fraction of the computational cost. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

3. Materials and Methods

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Taken together, these findings argue for treating the two regimes separately in future analyses. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We release the processing scripts and intermediate artefacts to support independent verification. Longer observation windows narrow the confidence interval but do not shift the central estimate. The measurements reported here were collected over twelve independent trials under controlled conditions. Sensitivity to the smoothing window proved modest across the range we examined.

We release the processing scripts and intermediate artefacts to support independent verification. Replicate agreement was high, suggesting the protocol is robust to small handling differences. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Sensitivity to the smoothing window proved modest across the range we examined. The reduced model retains most of the explanatory power at a fraction of the computational cost. Longer observation windows narrow the confidence interval but do not shift the central estimate.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. Neither of the alternative formulations improved on the simple additive model in our setting. Longer observation windows

narrow the confidence interval but do not shift the central estimate. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Taken together, these findings argue for treating the two regimes separately in future analyses. Calibration drift over the trial period remained below the threshold set out in the protocol. The measurements reported here were collected over twelve independent trials under controlled conditions.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The reduced model retains most of the explanatory power at a fraction of the computational cost. The measurements reported here were collected over twelve independent trials under controlled conditions. We release the processing scripts and intermediate artefacts to support independent verification. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance.

4. Experimental Setup

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Calibration drift over the trial period remained below the threshold set out in the protocol. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Sensitivity to the smoothing window proved modest across the range we examined. We release the processing scripts and intermediate artefacts to support independent verification. Taken together, these findings argue for treating the two regimes separately in future analyses.

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Neither of the alternative formulations improved on the simple additive model in our setting. The measurements reported here were collected over twelve independent trials under controlled conditions. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

5. Results

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Residual error grows approximately linearly with the number of coupled parameters. Neither of the alternative formulations improved on the simple additive model in our setting. The reduced model retains most of the explanatory power at a fraction of the computational cost. We release the processing scripts and intermediate artefacts to support independent verification. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Taken together, these findings argue for treating the two regimes separately in future analyses. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

Neither of the alternative formulations improved on the simple additive model in our setting. Sensitivity to the smoothing window proved modest across the range we examined. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Taken together, these findings argue for treating the two regimes separately in future analyses. Calibration drift over the trial period remained below the threshold set out in the protocol.

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Replicate agreement was high, suggesting the protocol is robust to small handling differences. We release the processing scripts and intermediate artefacts to support independent verification. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. Calibration drift over the trial period remained below the threshold set out in the protocol. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The reduced model retains most of the explanatory power at a fraction of the computational cost. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Longer observation windows narrow the confidence interval but do not shift the central estimate. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The measurements reported here were collected over twelve independent trials under controlled conditions.

Sensitivity to the smoothing window proved modest across the range we examined. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The reduced model retains most of the explanatory power at a fraction of the computational cost. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Calibration drift over the trial period remained below the threshold set out in the protocol. We release the processing scripts and intermediate artefacts to support independent verification. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Longer observation windows narrow the confidence interval but do not shift the central estimate.

6. Discussion

Neither of the alternative formulations improved on the simple additive model in our setting. Taken together, these findings argue for treating the two regimes separately in future analyses. The reduced model retains most of the explanatory power at a fraction of the computational cost. Residual error grows approximately linearly with the number of coupled parameters. We release the processing scripts and intermediate artefacts to support independent verification. The measurements reported here were collected over twelve independent trials under controlled conditions. Longer observation windows narrow the confidence interval but do not shift the central estimate. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Sensitivity to the smoothing window proved modest across the range we examined.

Neither of the alternative formulations improved on the simple additive model in our setting. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Sensitivity to the smoothing window proved modest across the range we examined. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The measurements reported here were collected over twelve independent trials under controlled conditions.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Longer observation windows narrow the confidence interval but do not shift the central estimate. Residual error grows approximately linearly with the number of coupled parameters. The measurements reported here were collected over twelve independent trials under controlled conditions.

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Longer observation windows narrow the confidence interval but do not shift the central estimate. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. The reduced model retains most of the explanatory power at a fraction of the computational cost. Sensitivity to the smoothing window proved modest across the range we examined. Residual error grows approximately linearly with the number of coupled parameters. We release the processing scripts and intermediate artefacts to support independent verification. Calibration drift over the trial period remained below the threshold set out in the

protocol. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

Neither of the alternative formulations improved on the simple additive model in our setting. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Calibration drift over the trial period remained below the threshold set out in the protocol. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The reduced model retains most of the explanatory power at a fraction of the computational cost.

7. Limitations

Sensitivity to the smoothing window proved modest across the range we examined. Calibration drift over the trial period remained below the threshold set out in the protocol. The reduced model retains most of the explanatory power at a fraction of the computational cost. We release the processing scripts and intermediate artefacts to support independent verification. Taken together, these findings argue for treating the two regimes separately in future analyses. Neither of the alternative formulations improved on the simple additive model in our setting.

Taken together, these findings argue for treating the two regimes separately in future analyses. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Sensitivity to the smoothing window proved modest across the range we examined. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance.

Taken together, these findings argue for treating the two regimes separately in future analyses. Calibration drift over the trial period remained below the threshold set out in the protocol. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. The reduced model retains most of the explanatory power at a fraction of the computational cost.

Calibration drift over the trial period remained below the threshold set out in the protocol. Longer observation windows narrow the confidence interval but do not shift the central estimate. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The reduced model retains most of the explanatory power at a fraction of the computational cost. Sensitivity to the smoothing window proved modest across the range we examined.

8. Conclusion

The reduced model retains most of the explanatory power at a fraction of the computational cost. We release the processing scripts and intermediate artefacts to support independent verification. Sensitivity to the smoothing window proved modest across the range we examined. The measurements reported here were collected over twelve independent trials under controlled conditions.

We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Longer observation windows narrow the confidence interval but do not shift the central estimate. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Taken together, these findings argue for

treating the two regimes separately in future analyses.

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The reduced model retains most of the explanatory power at a fraction of the computational cost. We release the processing scripts and intermediate artefacts to support independent verification. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

Sensitivity to the smoothing window proved modest across the range we examined. Taken together, these findings argue for treating the two regimes separately in future analyses. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. The measurements reported here were collected over twelve independent trials under controlled conditions. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Neither of the alternative formulations improved on the simple additive model in our setting. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Longer observation windows narrow the confidence interval but do not shift the central estimate.