

# Query Planners Revisited: A Latency-Aware Perspective

Hyun-woo Dusseault, Lucia Abernathy, Beatriz Fontenot

*Meridian Institute of Technology*

**Section:** Computing  
**Published:** 2024-12-20  
**DOI:** 10.5281/mjas.13.4.01  
**Keywords:** sensitivity analysis; reproducibility; field trial

## Abstract

We examine the behaviour of cache hierarchies in append-only stores. Calibration drift over the trial period remained below the threshold set out in the protocol. Sensitivity to the smoothing window proved modest across the range we examined. Residual error grows approximately linearly with the number of coupled parameters.

## 1. Introduction

Residual error grows approximately linearly with the number of coupled parameters. Sensitivity to the smoothing window proved modest across the range we examined. Calibration drift over the trial period remained below the threshold set out in the protocol. We release the processing scripts and intermediate artefacts to support independent verification.

The reduced model retains most of the explanatory power at a fraction of the computational cost. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Sensitivity to the smoothing window proved modest across the range we examined. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

## 2. Related Work

The measurements reported here were collected over twelve independent trials under controlled conditions. Taken together, these findings argue for treating the two regimes separately in future analyses. Sensitivity to the smoothing window proved modest across the range we examined. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Neither of the alternative formulations improved on the simple additive model in our setting.

We release the processing scripts and intermediate artefacts to support independent verification. Neither of the alternative formulations improved on the simple additive model in our setting. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. The measurements reported here were collected over twelve independent trials under controlled conditions. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Taken together, these findings argue for treating the two regimes separately in future analyses. Longer observation windows narrow the confidence interval but do not shift the

central estimate. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Sensitivity to the smoothing window proved modest across the range we examined.

Sensitivity to the smoothing window proved modest across the range we examined. We release the processing scripts and intermediate artefacts to support independent verification. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The reduced model retains most of the explanatory power at a fraction of the computational cost. Longer observation windows narrow the confidence interval but do not shift the central estimate. Residual error grows approximately linearly with the number of coupled parameters. The measurements reported here were collected over twelve independent trials under controlled conditions. Replicate agreement was high, suggesting the protocol is robust to small handling differences. Calibration drift over the trial period remained below the threshold set out in the protocol.

### **3. Materials and Methods**

Neither of the alternative formulations improved on the simple additive model in our setting. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The reduced model retains most of the explanatory power at a fraction of the computational cost. Calibration drift over the trial period remained below the threshold set out in the protocol. Sensitivity to the smoothing window proved modest across the range we examined. We release the processing scripts and intermediate artefacts to support independent verification. Residual error grows approximately linearly with the number of coupled parameters. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. The measurements reported here were collected over twelve independent trials under controlled conditions. Longer observation windows narrow the confidence interval but do not shift the central estimate. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. The reduced model retains most of the explanatory power at a fraction of the computational cost. Taken together, these findings argue for treating the two regimes separately in future analyses. Sensitivity to the smoothing window proved modest across the range we examined. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance.

Calibration drift over the trial period remained below the threshold set out in the protocol. Residual error grows approximately linearly with the number of coupled parameters. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The measurements reported here were collected over twelve independent trials under controlled conditions. We release the processing scripts and intermediate artefacts to support independent verification. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Longer observation windows narrow the confidence interval but do not shift the central estimate.

## 4. Experimental Setup

Calibration drift over the trial period remained below the threshold set out in the protocol. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Residual error grows approximately linearly with the number of coupled parameters. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

Residual error grows approximately linearly with the number of coupled parameters. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Taken together, these findings argue for treating the two regimes separately in future analyses. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.