

Characterizing Graph Embeddings in Multi-Tenant Pipelines

Wren Ibsen, Oleksandr Jovanovic, Dilnoza Abernathy

Calder Bay University

Section: Computing

Published: 2025-03-28

DOI: 10.5281/mjas.14.1.06

Keywords: cost accounting; finite elements; noise robustness

Abstract

This paper reports a systematic evaluation of federated scheduling heuristics. Residual error grows approximately linearly with the number of coupled parameters. The reduced model retains most of the explanatory power at a fraction of the computational cost. Calibration drift over the trial period remained below the threshold set out in the protocol.

1. Introduction

A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The measurements reported here were collected over twelve independent trials under controlled conditions. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Residual error grows approximately linearly with the number of coupled parameters. We release the processing scripts and intermediate artefacts to support independent verification. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

Residual error grows approximately linearly with the number of coupled parameters. Calibration drift over the trial period remained below the threshold set out in the protocol. Taken together, these findings argue for treating the two regimes separately in future analyses. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The reduced model retains most of the explanatory power at a fraction of the computational cost.

The measurements reported here were collected over twelve independent trials under controlled conditions. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Sensitivity to the smoothing window proved modest across the range we examined. We release the processing scripts and intermediate artefacts to support independent verification. Calibration drift over the trial period remained below the threshold set out in the protocol. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Neither of the alternative formulations improved on the simple additive model in our setting.

Neither of the alternative formulations improved on the simple additive model in our setting. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. Longer observation windows narrow the confidence interval but do not shift the central estimate. Our baseline reproduces the behaviour

described in earlier work to within the stated tolerance. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Residual error grows approximately linearly with the number of coupled parameters.

2. Related Work

Replicate agreement was high, suggesting the protocol is robust to small handling differences. The reduced model retains most of the explanatory power at a fraction of the computational cost. Neither of the alternative formulations improved on the simple additive model in our setting. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Taken together, these findings argue for treating the two regimes separately in future analyses.

Sensitivity to the smoothing window proved modest across the range we examined. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Residual error grows approximately linearly with the number of coupled parameters. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Longer observation windows narrow the confidence interval but do not shift the central estimate.

Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We release the processing scripts and intermediate artefacts to support independent verification. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Longer observation windows narrow the confidence interval but do not shift the central estimate. Calibration drift over the trial period remained below the threshold set out in the protocol. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance.

Residual error grows approximately linearly with the number of coupled parameters. Replicate agreement was high, suggesting the protocol is robust to small handling differences. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. We release the processing scripts and intermediate artefacts to support independent verification. The measurements reported here were collected over twelve independent trials under controlled conditions.

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Sensitivity to the smoothing window proved modest across the range we examined. Taken together, these findings argue for treating the two regimes separately in future analyses. Replicate agreement was high, suggesting the protocol is robust to small handling differences.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. Neither of the alternative formulations improved on the simple additive model in our setting. Sensitivity to the smoothing window proved modest across the range we examined. Residual error grows approximately linearly with the number of coupled parameters. Calibration drift over the trial period remained below the threshold set out in the protocol. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. We release the processing scripts and intermediate artefacts to support independent verification.

3. Materials and Methods

Residual error grows approximately linearly with the number of coupled parameters. Taken together, these findings argue for treating the two regimes separately in future analyses. We release the processing scripts and intermediate artefacts to support independent verification. Longer observation windows narrow the confidence interval but do not shift the central estimate. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. The measurements reported here were collected over twelve independent trials under controlled conditions. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The reduced model retains most of the explanatory power at a fraction of the computational cost. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

Residual error grows approximately linearly with the number of coupled parameters. Calibration drift over the trial period remained below the threshold set out in the protocol. We release the processing scripts and intermediate artefacts to support independent verification. The reduced model retains most of the explanatory power at a fraction of the computational cost. Neither of the alternative formulations improved on the simple additive model in our setting. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study.

Taken together, these findings argue for treating the two regimes separately in future analyses. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The measurements reported here were collected over twelve independent trials under controlled conditions. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Sensitivity to the smoothing window proved modest across the range we examined. Longer observation windows narrow the confidence interval but do not shift the central estimate.

4. Experimental Setup

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Longer observation windows narrow the confidence interval but do not shift the central estimate. Neither of the alternative formulations improved on the simple additive model in our setting. The reduced model retains most of the explanatory power at a fraction of the computational cost. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Sensitivity to the smoothing window proved modest across the range we examined. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. The measurements reported here were collected over twelve independent trials under controlled conditions.

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Calibration drift over the trial period remained below the threshold set out in the protocol. Longer observation windows narrow the confidence interval but do not shift the central estimate. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. Neither of the alternative formulations improved on the simple additive model in our setting. The measurements reported here were collected over twelve independent trials under controlled conditions. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Taken together,

these findings argue for treating the two regimes separately in future analyses. Neither of the alternative formulations improved on the simple additive model in our setting. Residual error grows approximately linearly with the number of coupled parameters.

5. Results

The reduced model retains most of the explanatory power at a fraction of the computational cost. The measurements reported here were collected over twelve independent trials under controlled conditions. Neither of the alternative formulations improved on the simple additive model in our setting. Replicate agreement was high, suggesting the protocol is robust to small handling differences. We release the processing scripts and intermediate artefacts to support independent verification.

Replicate agreement was high, suggesting the protocol is robust to small handling differences. Longer observation windows narrow the confidence interval but do not shift the central estimate. Our baseline reproduces the behaviour described in earlier work to within the stated tolerance. Calibration drift over the trial period remained below the threshold set out in the protocol. Discrepancies at the boundary are consistent with the simplifying assumptions noted above.

Taken together, these findings argue for treating the two regimes separately in future analyses. Longer observation windows narrow the confidence interval but do not shift the central estimate. Neither of the alternative formulations improved on the simple additive model in our setting. Replicate agreement was high, suggesting the protocol is robust to small handling differences. The measurements reported here were collected over twelve independent trials under controlled conditions. We release the processing scripts and intermediate artefacts to support independent verification. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Sensitivity to the smoothing window proved modest across the range we examined.

Neither of the alternative formulations improved on the simple additive model in our setting. Calibration drift over the trial period remained below the threshold set out in the protocol. The reduced model retains most of the explanatory power at a fraction of the computational cost. We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity.

We observe a pronounced inflection once the load exceeds roughly sixty percent of nominal capacity. A modest but repeatable bias appears in the lowest-magnitude bin and warrants further study. Discrepancies at the boundary are consistent with the simplifying assumptions noted above. We release the processing scripts and intermediate artefacts to support independent verification. Neither of the alternative formulations improved on the simple additive model in our setting. Taken together, these findings argue for treating the two regimes separately in future analyses. Longer observation windows narrow the confidence interval but do not shift the central estimate.